



# 703308 VO High-Performance Computing WS2021/2022

## Debugging Parallel Programs

Philipp Gschwandtner

# Overview

---

- ▶ functional debugging

- ▶ generic guidelines
- ▶ serial debugging
- ▶ parallelism-specific debugging

- ▶ performance debugging

- ▶ generic guidelines
- ▶ serial debugging
- ▶ parallelism-specific debugging

# Motivation

---



[https://www.youtube.com/watch?v=gp\\_D8r-2hwk](https://www.youtube.com/watch?v=gp_D8r-2hwk)

# Motivation

---

- ▶ Why do we need debugging?
  - ▶ Because we make mistakes!
- ▶ Why do we need a lecture about this?
  - ▶ OpenMPI FAQ “Debugging applications in parallel”, first question:  
Q: “How do I debug OpenMPI processes in parallel?”  
A: “This is a difficult question. [...] This FAQ section does not provide any definite solutions to debugging in parallel. [...]”



# Functional Debugging



# Functional Debugging

---

- ▶ everything that results in not getting the correct program output
  - ▶ program crashes
  - ▶ program not finishing (freezes, infinite loops)
  - ▶ incorrect output
- ▶ errors can be deterministic or non-deterministic
  - ▶ ensure/maximize reproducibility during testing (e.g. fix random seeds, process/thread numbers, affinities, ...)
- ▶ all that applies to debugging serial programs is crucial for parallel ones
  - ▶ If you can't trust the serial implementation, why would you in a parallel context?

# Coding Guidelines

---

- ▶ write clean code that prevents bugs or facilitates their detection, e.g.
  - ▶ use meaningful identifiers
  - ▶ minimize vertical distance of variables
  - ▶ don't use OpenMP's `private`
  - ▶ follow the Don't Repeat Yourself (DRY) principle (single component per feature)
  - ▶ ...
- ▶ The toolchain you must use!
  - ▶ read & heed compiler warnings
  - ▶ write and regularly run unit and/or integration tests, especially aimed at (varying degrees of) parallelism
  - ▶ use code coverage tests
  - ▶ use continuous integration
  - ▶ use source version control



# “Best of” Real Commit Messages Encountered in the Past

---


- ▶ stuff
- ▶ manager stuff
- ▶ more manager stuff
- ▶ Make things work
- ▶ ::w  
:q  
Merge branch 'master'
- ▶ dl;adlwa
- ▶ Added performance fix for DataItemManager::get() by caching fragment result in reference
- ▶ Removed debug print statement
- ▶ Fixed a linking issue of the unwrap\_tuple function
- ▶ Redirected runtime system output to error stream
- ▶ fixing typos



## Recommended Reading/Reference Material

---

- ▶ “Clean Code” by Robert Martin, Prentice Hall 2008
  - ▶ ISBN 9780132350884
  - ▶ also available in German
- ▶ naming, functions, commenting, formatting, data structures, error handling, unit tests, classes, concurrency, refinement & refactoring, ...



# Generic Debugging Guidelines

---

- ▶ create a **Minimal Working Example (MWE)**
  - ▶ minimize problem size
  - ▶ minimize software components/features involved
  - ▶ ensure/increase reproducibility
  - ▶ if parallel
    - ▶ minimize machine size (number of threads and/or ranks)
    - ▶ minimize complexity of parallel interaction (e.g. communication patterns, ...)
- ▶ minimizes debugging feedback cycles times, amount of memory to inspect, amount of code to consider, overall degree of complexity of component & parallel interaction
  - ▶ sounds simple, but don't underestimate this
  - ▶ every change along the way to an MWE gives you more information about the problem

# Serial Debuggers

---

## ▶ gdb

- ▶ useful for inspecting memory contents and getting call stacks
- ▶ can work with multi-threaded programs and also MPI
  - ▶ `mpiexec -n X gdb -ex 'run' -ex 'bt' -ex 'quit' ./a.out`
- ▶ can be used to debug a single MPI process among many
  - ▶ `mpiexec -n 1 gdb ./a.out : -n X-1 ./a.out`
- ▶ can be attached to already-running processes
  - ▶ `gdb -pid 12345`

## ▶ valgrind

- ▶ mostly used for finding memory leaks (can also simulate cache or generate call graph)
- ▶ can work with multi-threaded programs (but no parallel execution!)
- ▶ can yield some false positives e.g. for OpenMP related to thread-local storage

## Sanitizers (Still Mostly Serial)

---

- ▶ tools that instrument code at compile time to perform checks at runtime
  - ▶ often lower overhead compared to external tools such as valgrind
  - ▶ if in doubt, check same issue with multiple tools (e.g. address sanitizers of multiple compilers and valgrind)
- ▶ depending on compiler, several sanitizers available, e.g.
  - ▶ address: buffer overflows, use-after-free, stack corruption, etc.
  - ▶ undefined behavior: signed integer overflow, float division by zero, negative shift operands, etc.
  - ▶ thread: detects data races
  - ▶ leak: detects memory leaks

# Call Graph Generators

---

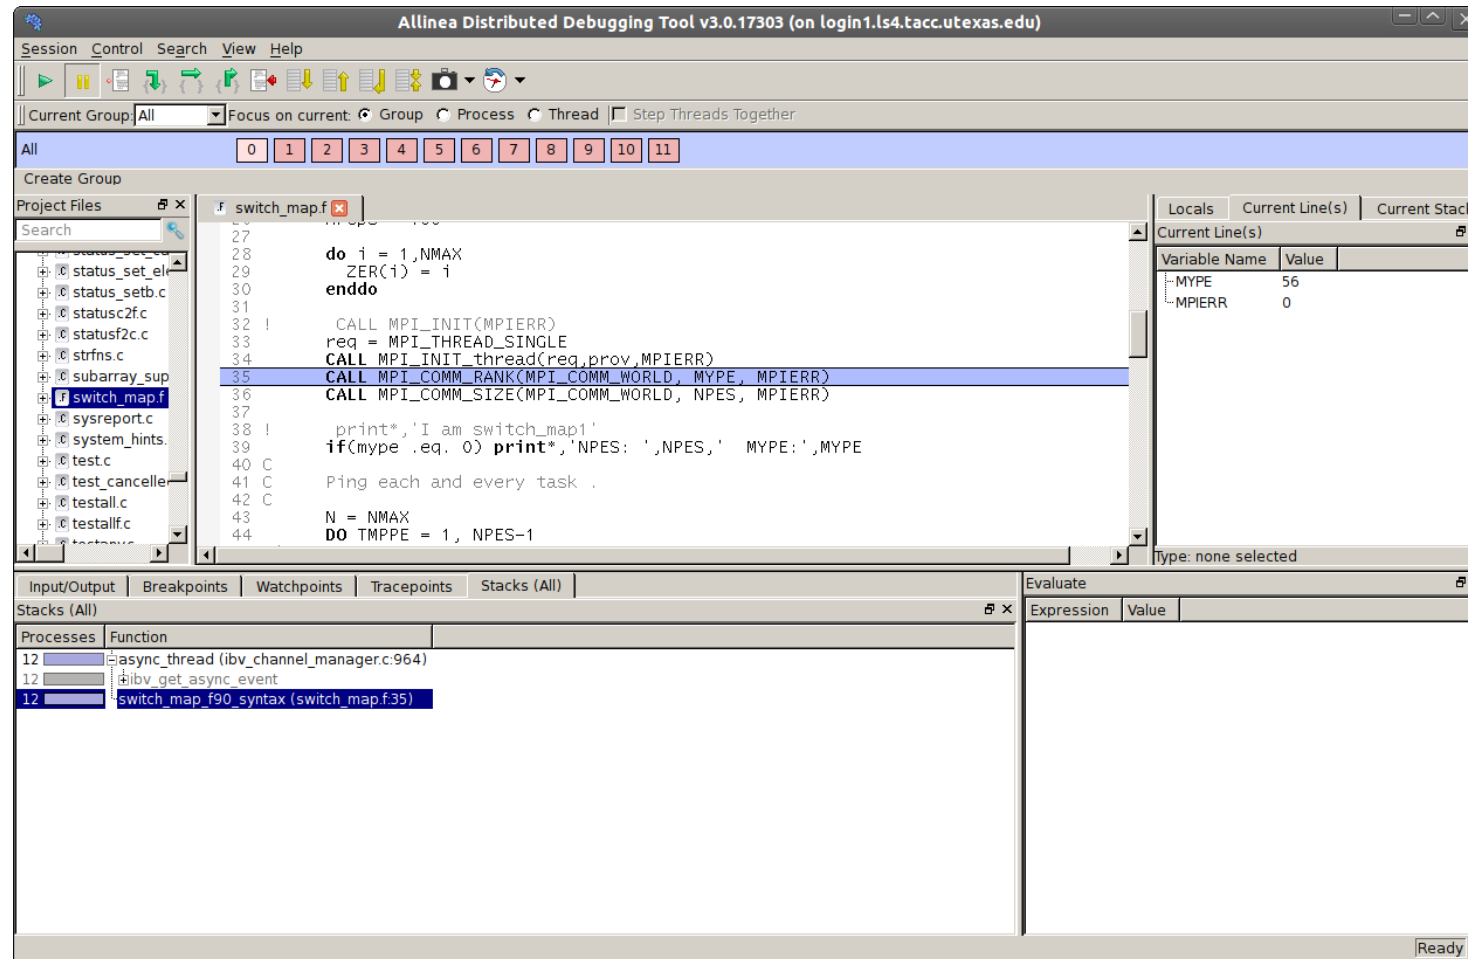
- ▶ many tools available for generating call graphs
- ▶ static (at compile time)
  - ▶ doxygen, opt (llvm), cflow (gcc), etc.
- ▶ dynamic (at runtime)
  - ▶ gprof, callgrind, OpenPAT, pprof, CodeAnalyst, etc.
  - ▶ most performance analysis tools offer some form of call graph generation

# Parallel Debuggers

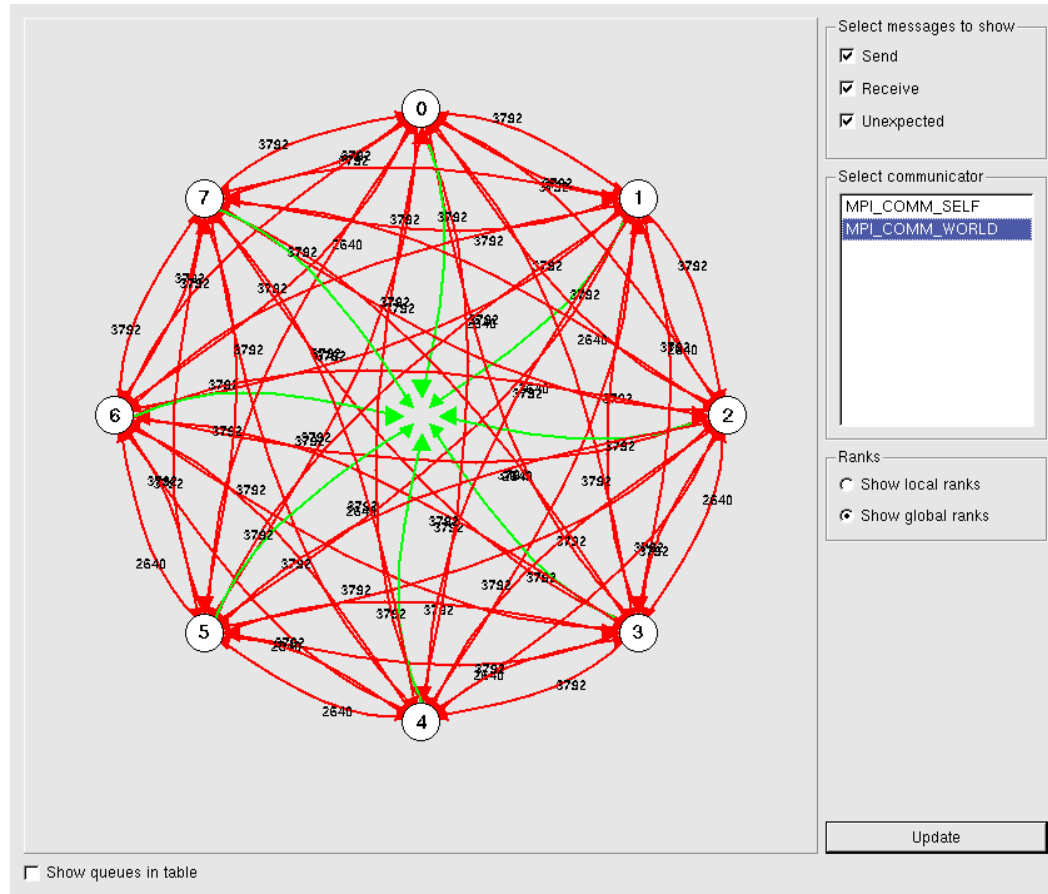
---

- ▶ very little free software
- ▶ two commercial top dogs: DDT (ARM) and TotalView (Rogue Wave Software)
- ▶ support OpenMP, MPI, CUDA, etc.
- ▶ several features centered around parallelism
  - ▶ examine variables per rank/thread, examine send/receive queues of MPI libraries, etc.
  - ▶ still, limited usefulness

# DDT Screenshot (Overview)



# DDT Screenshots (Communication Patterns, Data Across Ranks)



Locals	
Name	Value
argc	1
argv	0x7ffffffdc58
beingWatched	0
bigArray	
dest	0
dynamicArray	0x818020
environ	0x7ffffffdea0
i	0
message	""
my_rank	0
p	512
source	32767
status	
t2	0x603050
tables	
tag	50
test	
x	10000
y	12



# Automatic Race Condition Debugging

---

- ▶ difficult to do automatically and exactly
  - ▶ statically detecting race conditions is NP-hard
  - ▶ dynamically detecting race conditions incurs large runtime overhead (every memory access and synchronization action must be logged and checked)
- ▶ most solutions resort to heuristics
  - ▶ several experimental tools available in research
  - ▶ many issues: limited scope, only apply to a subset of programming language, etc.
  - ▶ few “mature” tools, e.g. Intel Inspector

# Intel Inspector

---

## ▶ features

- ▶ free
- ▶ Linux & Windows version
- ▶ automatically finds bugs in multi-threaded programs
  - ▶ deadlocks
  - ▶ memory corruption
  - ▶ race conditions
  - ▶ vulnerabilities
- ▶ supports OpenMP, TBB, Pthreads, Windows threads

## ▶ limitations & issues

- ▶ slowdown by 1-2 orders of magnitude!
- ▶ explicit support only for Intel OpenMP runtime
- ▶ error detection only at runtime, only in executed control flow branches
- ▶ false positives and negatives possible

# OpenMP Data Race Example 1

---

```
int counter = 0;

#pragma omp parallel for
for(int i = 0; i < 10; ++i) {
    counter++;
}
```

Description ▲	Source	Function	Module
Read	ConsoleApplication1.cpp:9	main	consoleapplication1.exe
7	#pragma omp parallel for		
8	for (int i = 0; i < 10; ++i) {		
9	counter++;		
10	}		
11			
Write	ConsoleApplication1.cpp:9	main	consoleapplication1.exe
7	#pragma omp parallel for		
8	for (int i = 0; i < 10; ++i) {		
9	counter++;		
10	}		
11			

## OpenMP Data Race Example 2

```
int sum = 0;

#pragma omp parallel for
for(int i = 0; i < 10; i++) {
    int tmp = sum;
    tmp = tmp + 1;
    sum = tmp;
}
```

Description ▲	Source	Function	Module
Read	ConsoleApplication1.cpp:17	main	consoleapplication1.exe
15	#pragma omp parallel for		
16	for (int i = 0; i < 10; i++) {		
17	int tmp = sum;		
18	tmp = tmp + 1;		
19	sum = tmp;		
Write	ConsoleApplication1.cpp:19	main	consoleapplication1.exe
17	int tmp = sum;		
18	tmp = tmp + 1;		
19	sum = tmp;		
20	}		
21			

## OpenMP Data Race Example 2: Wrong Fix

---

```
int sum = 0;

#pragma omp parallel for
for(int i = 0; i < 10; i++) {
    int tmp;
    #pragma omp critical
    tmp = sum;
    tmp = tmp + 1;
    #pragma omp critical
    sum = tmp;
}
```





(not detected by Intel Inspector 2020)

# Domain-specific Debugging

---

- ▶ Visualize the output using appropriate tools
  - ▶ gnuplot
  - ▶ ParaView
  - ▶ ...
- ▶ note that this usually prohibits automatic checking
  - ▶ whenever feasible, unit and integration tests are preferred



# Best Approach to Debugging Parallel Programs

---

- ▶ know your algorithm and implementation
  - ▶ e.g. “an n-body simulation using Barnes-Hut”
- ▶ know your programming models and languages, and their semantics
  - ▶ “OpenMP threadprivates persist per thread between parallel regions with the same number of threads and affinity policies”
  - ▶ “this C++ object’s destructor will be called at the end of the full-expression”
- ▶ Don’t trust (seemingly) automatic analysis tools too much, read and understand the source code when available!



# Performance Debugging





# Performance Debugging

---

- ▶ also sometimes known as “*non-functional*” debugging (not related to functional output)
  - ▶ short execution time not necessarily but most often the only goal
  - ▶ much more tricky than functional debugging
    - ▶ How do you know the performance bug was fixed? Because it’s “faster” now?
- ▶ most aspects of functional debugging or sequential programs still apply
  - ▶ coding guidelines & best practice
  - ▶ + reproducibility (e.g. fix random seeds, scheduling affinities, ...)
  - ▶ + if required, performance unit tests, performance regression checks
  - ▶ + performance tools (the ones for sequential programs can also be useful)
  - ▶ + a lot more knowledge about hardware required

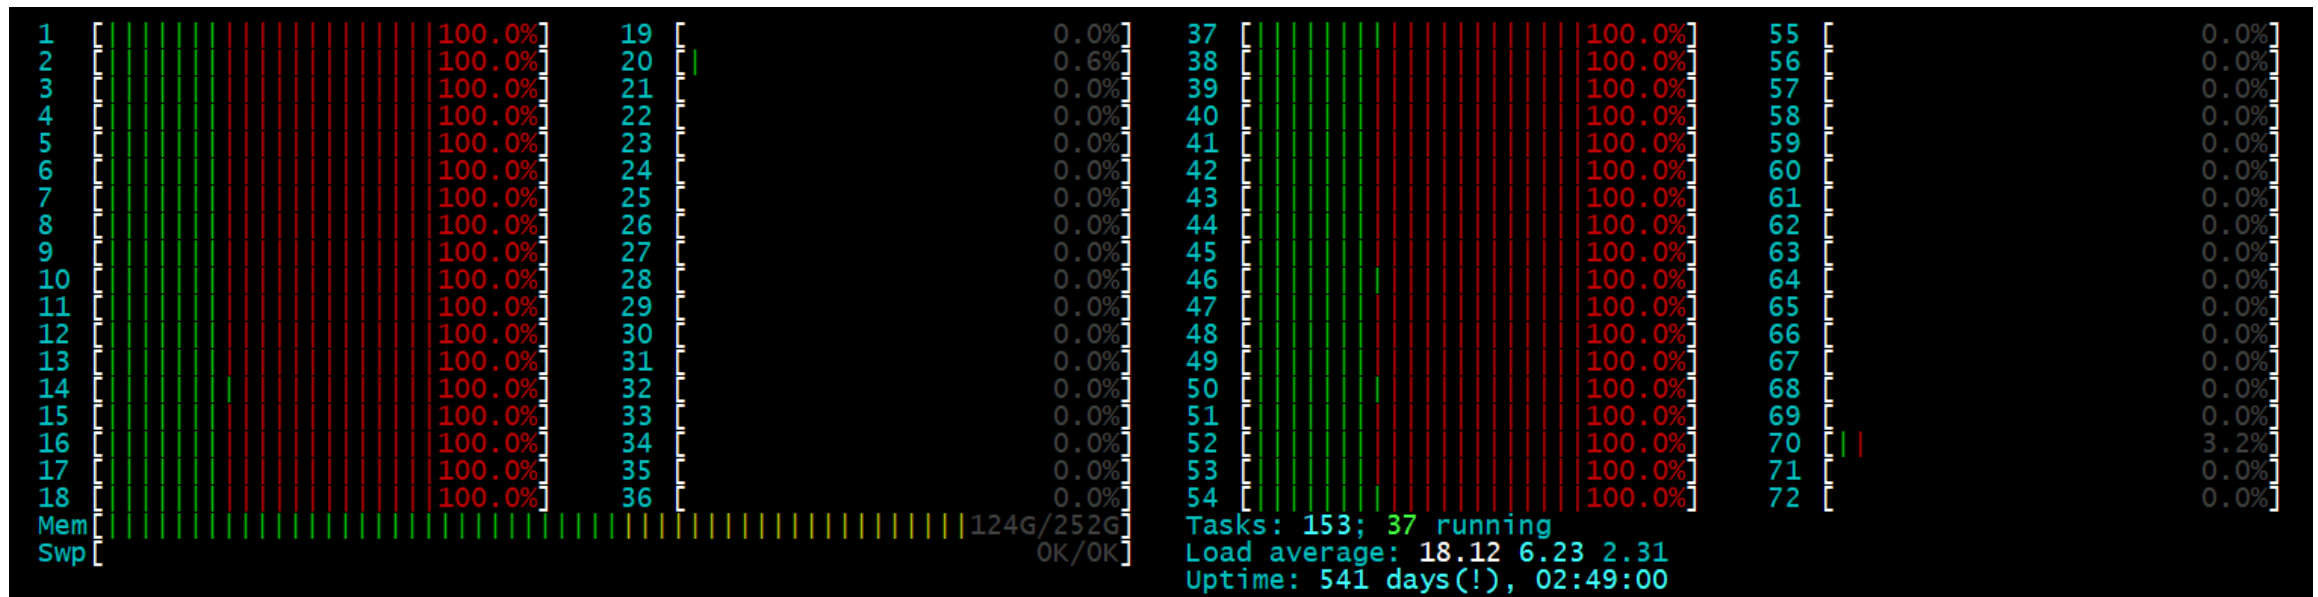
# (h)top

---

- ▶ Don't underestimate the power of top or htop!
- ▶ Get a high-level overview of the workload on the system (and it's components) and compare to what you expected!
  - ▶ What's the ratio between user time and system time?
    - ▶ high system time could be caused by inefficient I/O, high amount of context switching, etc.
  - ▶ Which CPU cores am I really using?
    - ▶ the only way to verify affinity policies
  - ▶ What is the actual memory footprint vs. what it should be?
    - ▶ detect existence of memory leaks without any additional analysis tools

# htop & affinity

- ▶ 2x Intel E5-2699 v3 (18 cores per CPU) in a single node
- ▶ htop shows cores 1-18 and 37-54 busy, hence 36 cores total – right?



## Recap: perf

---

```
[c703429@login.lcc2 ~]$ perf stat ./heat_stencil_1D_seq
...
28,826,239,136 cycles:u          #    2.471 GHz
35,220,856,783 instructions:u   #    1.22  insn per cycle
 6,711,849,029 branches:u       # 575.356 M/sec
    1,295,209 branch-misses:u   #    0.02% of all branches
      1,044 LLC-load-misses:u
        26 LLC-store-misses:u
    15,312,122 L1-dcache-load-misses:u
    476,440,489 L1-dcache-store-misses:u
```

# Terminology

---

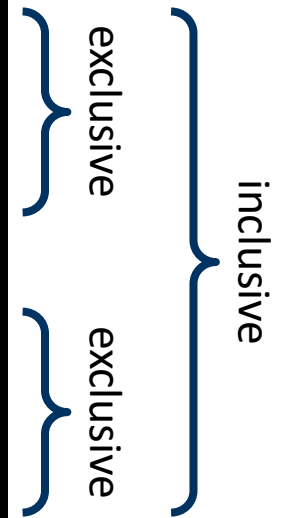
## ▶ instrumentation

- ▶ add source/machine code that will measure something when executed
- ▶ can happen manually, automatically, during compilation, linking, runtime, ...
- ▶ do not confuse with “measurement”

## ▶ inclusive/exclusive measurements

- ▶ do measurements include data for nested code regions (e.g. functions)?

```
int outside() {  
    for(int i = 0; i < N; ++i) {  
        // work  
    }  
    inside();  
    for(int j = 0; j < M; ++j) {  
        // more work  
    }  
}
```



# More Terminology: Sample- vs. Trace-based Profiling

---

## ▶ Sampling

- ▶ gives aggregated information of how much time spent where in the code
- ▶ based on statistics: does not provide information on the order of events, their time interval or exact numbers
- ▶ easy to accomplish, comparatively low overhead, no code changes required
  - ▶ stop program periodically and read program counter of CPU
  - ▶ build histogram at the end

## ▶ Tracing

- ▶ produces a detailed log of which event happened at what point in time
- ▶ allows to establish order of events, even across processes/nodes if clocks are in sync
- ▶ requires code changes/instrumentation
  - ▶ e.g. wrap every function call with  
`start_timer();`  
`func_call();`  
`end_timer();`

# gprof

---

- ▶ sample-based profiler
  - ▶ also limited code instrumenter for call graph generation and call counts
  - ▶ very simplistic, not always accurate
- ▶ available with every GCC installation
- ▶ very simple in its use
  - ▶ compile with debug symbols (-g) and gprof support (-pg)
  - ▶ run binary as usual
  - ▶ run `gprof binary gmon.out` to view results
  - ▶ use `--line` to get more detailed, line-based results

# gprof Example

---

```
int foo() {  
    long long counter = 0;  
    #pragma omp parallel for  
    for(int i = 0; i < N; ++i) {  
        #pragma omp critical  
        counter++;  
    }  
    return counter;  
}
```

```
int bar() {  
    long long partSum[MAX_NUM_THREADS][8];  
    long long counter = 0;  
    #pragma omp parallel  
    {  
        int tid = omp_get_thread_num();  
        partSum[tid][0] = 0;  
        #pragma omp for  
        for(int i=0; i<N; ++i) partSum[tid][0]++;  
        #pragma omp critical  
        counter += partSum[tid][0];  
    }  
    return counter;  
}
```



## gprof Example cont'd

---

Flat profile:

Each sample counts as 0.01 seconds.


% time	cumulative seconds	self seconds	calls	self Ts/call	total Ts/call	name
100.71	0.02	0.02				foo._omp_fn.0 (test.c:13 @ 400a3d)
0.00	0.02	0.00	1	0.00	0.00	bar (test.c:19 @ 40092c)
0.00	0.02	0.00	1	0.00	0.00	foo (test.c:8 @ 4008e6)

# gperftools

---

- ▶ sample-based profiler
  - ▶ formerly Google Performance Tools
- ▶ actually a collection of performance analysis tools and high-performance multi-threaded memory allocators
- ▶ very simple in its use
  - ▶ install gperftools library
  - ▶ link with `-lprofiler`
  - ▶ run with environment variable `CPUPROFILE=prof.out`
  - ▶ run `pprof binary prof.out` to view results (`--gv` for graphical visualization)

# gperftools Example


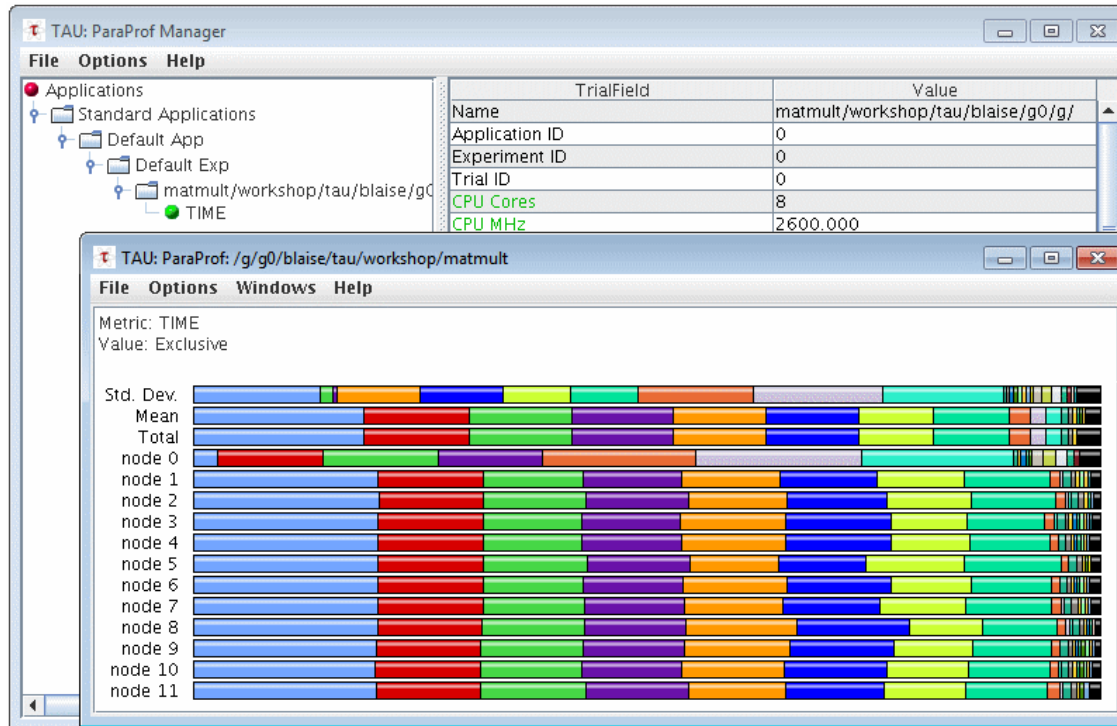


# Performance Analysis Tools for Parallel Programs

---

- ▶ **profiling and analysis software**
  - ▶ Intel Pin: dynamic binary instrumentation
  - ▶ Intel VTune: performance analysis for multi-threaded programs
  - ▶ Intel Advisor: dependency, vectorization and cache analysis tool
  - ▶ AMD CodeXL / NVIDIA Nsight: profiler and debugger for GPUs
  - ▶ TAU: profiling and tracing toolkit
  - ▶ PAPI: library for access to hardware performance counters
  - ▶ OProfile: sampling-based profiler with hardware performance counter support
  - ▶ also, some software built into your IDE, e.g. MS Visual Studio
- ▶ **analysis and visualization/reporting tools**
  - ▶ Scalasca, Vampir, Paraver, JumpShot, paraprof, CUBE, etc.
- ▶ These lists are by far not complete!

# TAU & ParaProf



# Tau Instrumentation Files

---


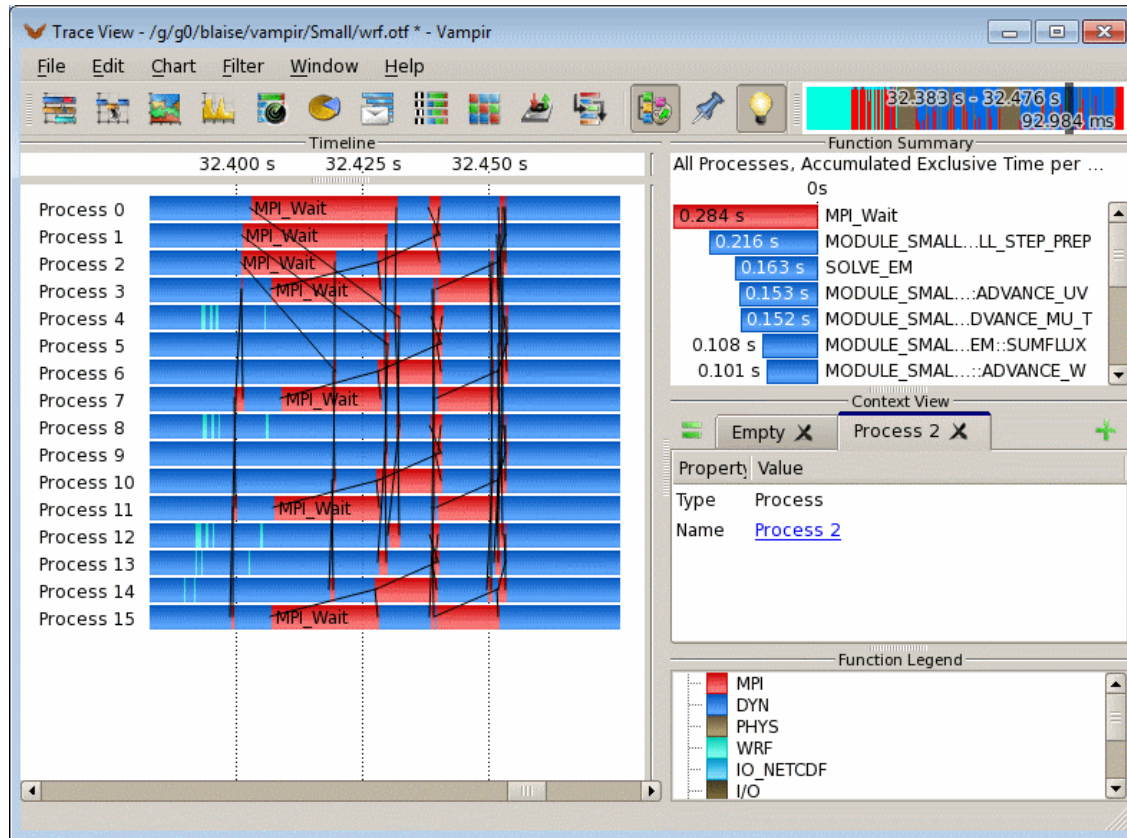
- ▶ allows to control the scope of instrumentation
  - ▶ reduces measurement overhead
  - ▶ reduces collected data to relevant content
- ▶ allows to select function patterns, loops, code lines, etc.

```
BEGIN_INSTRUMENT_SECTION
```

```
loops file="foo.cpp" routine="int  
    bar(int*, double)"
```

```
END_INSTRUMENT_SECTION
```

# Vampir



# General Hints When Working With Debuggers


---

- ▶ `-g` when compiling if source locations are required
- ▶ careful with optimization flags, especially `-O#`
  - ▶ function inlining, loop fusion/fission, ...
  - ▶ likely to obfuscate source code locations
  - ▶ if feasible, work in `-O0` or temporarily disable conflicting flags
- ▶ check whether child processes are included in analysis/reports
- ▶ check whether threads are included in analysis/reports
- ▶ if tracing or otherwise large-overhead instrumentation required, restrict to code regions of interest



# Points of Attack in Order of Benefit

---



# Summary

---

- ▶ functional debugging

- ▶ adhere to coding guideline and best practices of software engineering
- ▶ especially relevant for parallelism: know your programming models and semantics, don't trust automatic tools blindly

- ▶ performance debugging

- ▶ don't underestimate the power of simple tools
- ▶ many more advanced tools out there, but not straight-forward to use
- ▶ know your hardware and your program hotspots

# Image Sources

---

- ▶ Yoda: <https://www.deviantart.com/biggiepoppa/art/Master-Yoda-Star-Wars-395511111>
- ▶ DDT: <https://portal.tacc.utexas.edu/software/ddt>, [https://www.sharcnet.ca/help/index.php/Parallel\\_Debugging\\_with\\_DDT](https://www.sharcnet.ca/help/index.php/Parallel_Debugging_with_DDT), <https://developer.arm.com/docs/101136/latest/ddt/viewing-variables-and-data>
- ▶ Domain-specific debugging: <https://twitter.com/maven2mars/status/984440044659159040>, <https://www.nasa.gov/ames/image-feature/nasa-highlights-simulations-at-supercomputing-conference-like-aircraft-landing-gear>, ZAMG Wettervorhersage 06.10.2020 12:00
- ▶ TAU & ParaProf: <https://hpc.llnl.gov/software/development-environment-software/tau-tuning-and-analysis-utilities>
- ▶ Vampir: <https://hpc.llnl.gov/software/development-environment-software/vampir-vampir-server>